

IDEA

Majlis

Multiple AI models. One moderated room. You steer; the Curator keeps it sharp.

Hackathon · AMD GPU · 2026

PROBLEM

- Single models feel like a black box — no tension, no second opinions.
 - Users need concise, comparable answers — not endless essays.
 - Trust should adapt: bad panelists get kicked; good ones earn reputation.
-

PRODUCT

- **Curator** analyzes the question and recommends a panel.
 - **Round-based chat** — your prompt, then parallel model cards (stack / flip).
 - **Curator moves** — ready prompts (challenge, consensus, devil's advocate...).
 - **Kick + feedback** — multi-select reasons; removed models stay out of discussion context.
 - **Tools** — Tavily web search (per model if enabled; optional shared “web this round”).
-

STACK

- **Frontend:** React, Vite, TypeScript, Tailwind, shadcn/ui, Clerk, Cloudflare Pages.
 - **Backend:** FastAPI, SQLAlchemy/SQLite, JWT (Clerk).
 - **Inference:** vLLM on AMD (ROCm) — multiple 8B-class models + small curator.
 - **Edge:** Nginx on droplet; Cloudflare proxy + DNS.
-

ARCHITECTURE

- Browser → Pages → API (HTTPS) → Nginx → FastAPI.
 - FastAPI streams SSE from `/discuss`; messages persisted per room.
 - Each model has its own OpenAI-compatible vLLM port; curator calls are sync/stream as needed.
-

DEMO PATH

- Sign in → ask a question → Curator picks models → start room.
 - Read stacked cards → Curator move → models respond to each other (active panel only).
 - Kick a weak answer → Curator notes suggestion; reputation updates backend-side.
-

THANKS

Open the Majlis.

mohessaid.com · api + pages deployment ready.